

What does censorship-resistance look like in `code`?

@glozow

6B00 2C6E A3F9 1B1B 0DF0 C9BC 8F61 7F12 00A6 D25C



censorship-resistance is one of Bitcoin's most valuable properties



anyone can pay anyone, anywhere in the world



no central authority that can confiscate/freeze your money



even in the presence of adversaries

... but what does that look like in code?



```
send_payment()
```



```
send_payment_uncensorable()
```

I won't display code, but am linking code for slides:

tinyurl.com/seoul-code-links

 <- look for this icon on slide

which payment is easier to stop?



001



011




which payment is easier to stop?

this one



our equivalent is a distributed network of nodes



... but what does that look like in code?



we can run a decentralized network across `us-east-1`, `eu-west-1` and `ap-northeast-1`

woah then nobody can shut it down

We want many independent users to run nodes. From a software engineering perspective, this necessitates:



the software must be supported on many platforms




the code is open source and free to use



node operators have control of what is running



**running a node must be cheap
i.e. minimal operating costs,
even as usage scales**



mining nodes accept
transactions to collect fees

non-mining nodes
don't get paid



The Costs

every node has limited:

- memory
- bandwidth
- computation



The Costs

every node has limited:

- memory
- bandwidth
- computation

tx relay requires:

- memory
- bandwidth
- computation



The Costs

every node has limited:

- memory
- bandwidth
- computation

tx relay requires:

- memory
- bandwidth
- computation

block validation requires:

- memory
- bandwidth
- computation




The Benefits

Transaction relay amortizes block validation:

- loading UTXOs from disk
- downloading block data
- verifying signatures and scripts

This means:

- faster block relay (steadily instead of in spikes)
 - when network-wide = fewer reorgs
- more accurate fee estimation



The challenge = write useful code that can run with constrained resources
(e.g. a raspberry pi or a laptop app)



transaction relay vs
memory



Use memory wisely

Several caches exist:

- UTXO (“coins”) cache
- signature cache
- script verification result cache
- memory pool of validated transactions
- orphan pool of transactions for which we are missing UTXOs
- ring buffer of rejected and replaced transactions



A mempool is just a cache

= a memory pool of transactions that may end up in a block

- 👉 memory bounds: limit maximum usage (300MB)
- 👉 hit rate: what ultimately matters (compact block reconstruction)
- 👉 define a utility metric and eviction policy



Pop Quiz:

What happens when there are so many transactions that your mempool exceeds 300MB?



utility metric = fees

a tx is useful if it ends up being confirmed

=> a tx is mined if it would be profitable

=> incentive compatibility* is a good utility metric

Eviction policy:

- remove transactions > 2 weeks old
- skim away transactions with the worst fees*
- allow replacements with better fees*

*: it's complicated



“free relay” and min relay fee

- 👉 fees are not paid to the node operator
- 👉 however, fees represent a cost to tx creator



free relay = when total cost < memory (+bandwidth) used at 1sat/vB

- **-minrelaytxfee**: all txns must pay 1sat/vB, even if mempool is empty
- RBF rule #4: replacement must pay additional fees that cover the new transaction's size at 1sat/vB
- the tx also needs to be valid (no ultimate cost if it never confirms!)




transaction relay vs
bandwidth usage



everybody gossip to everybody:

- very fast propagation!
- $O(n^2)$ tx messages
- lots of wasted bandwidth
- they might even be invalid



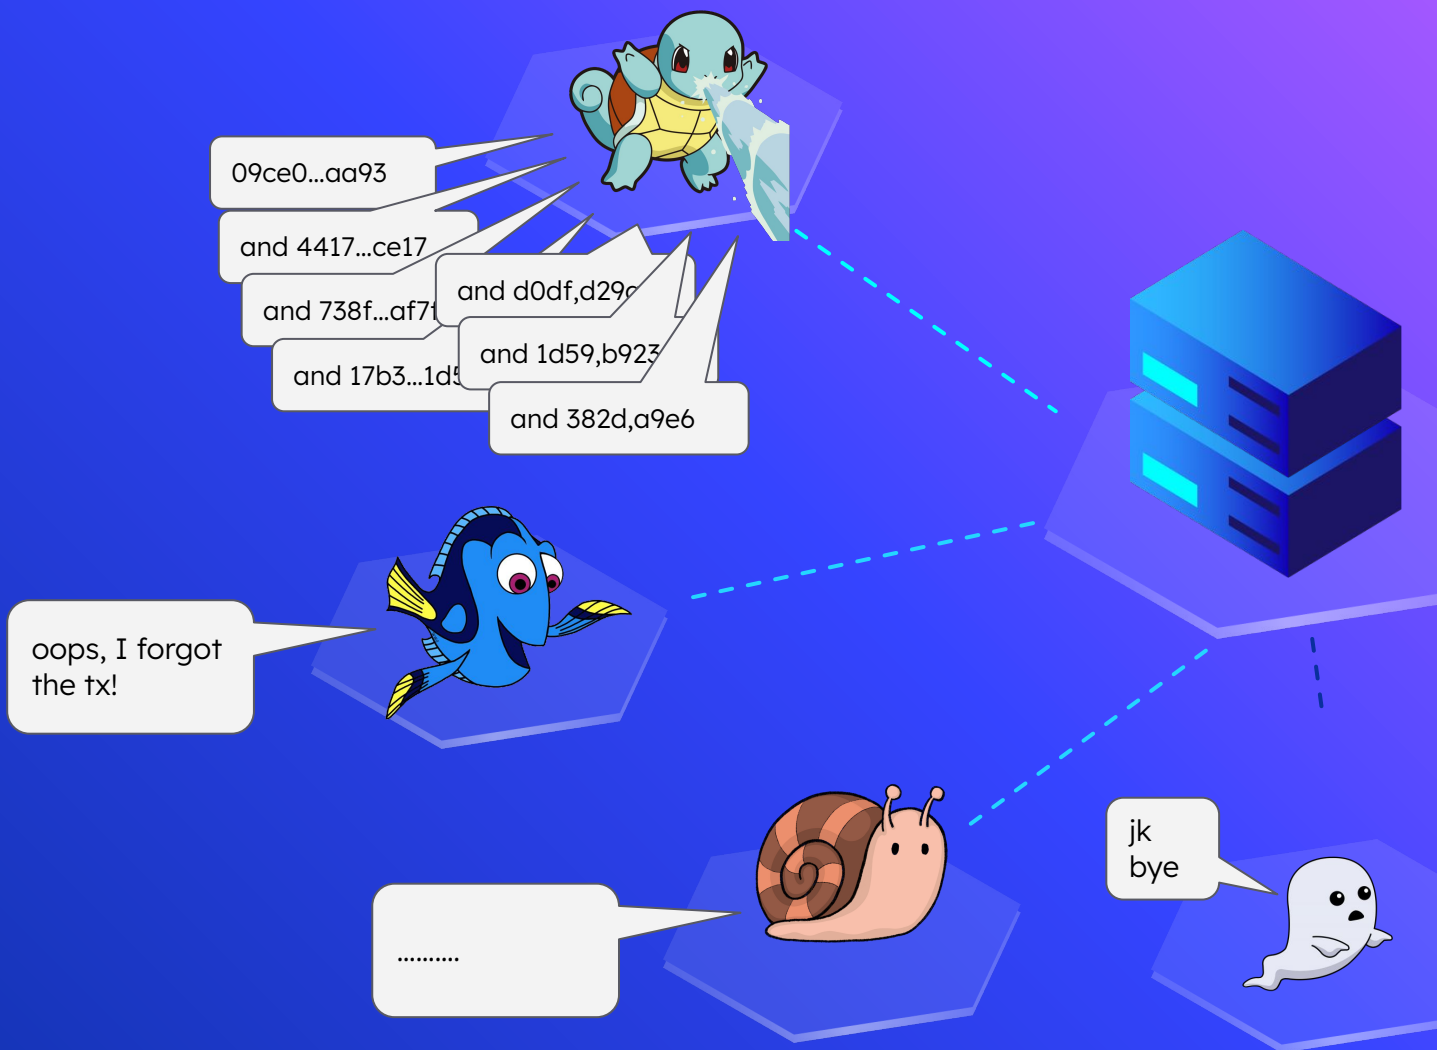
instead:

- gossip inventory by tx hash
- request what is missing from 1 peer
 - $O(n^2)$ 32b invs
 - $O(n)$ tx messages, usually
- don't re-download rejected transactions



What if they:

- flood
- forget
- stall
- drop off



We need to track all announcers

- double up requests after some timeout
- have preferences (e.g. outbounds)
- load-balance between peers



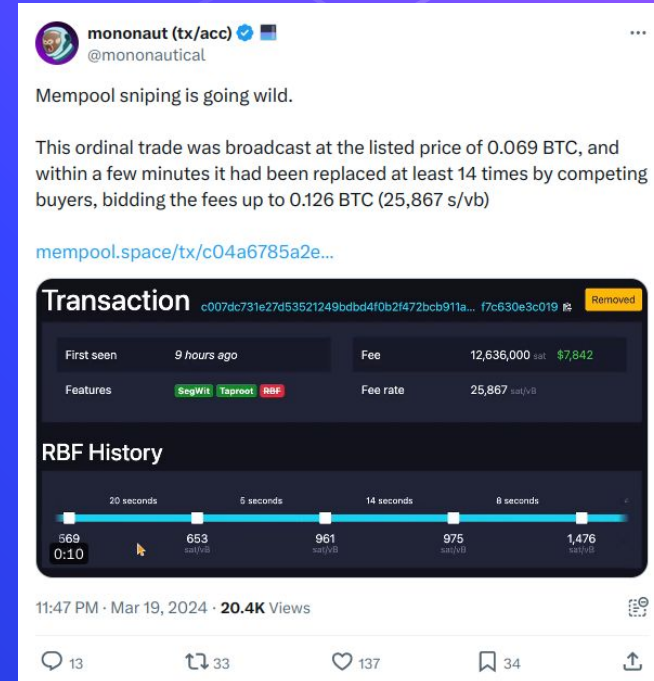
How else can somebody sabotage a tx?

- malleation of tx data
- pinning through attaching children or replacements



“mempool sniping”

- (1) seller creates a malleable transaction, signed SIGHASH_SINGLE | SIGHASH_ANYONECANPAY, exchanging the NFT for the sell price
- (2) buyer signs PSBT providing funds and sending NFT to themselves
- (3) anyone can malleate the tx after broadcast, providing new funds and changing NFT destination (original buyer does not lose their money)



The screenshot shows a tweet from user **mononaut (tx/acc)** (@mononautical) with the text "Mempool sniping is going wild." Below the text, it says: "This ordinal trade was broadcast at the listed price of 0.069 BTC, and within a few minutes it had been replaced at least 14 times by competing buyers, bidding the fees up to 0.126 BTC (25,867 s/vb)". A link to mempool.space/tx/c04a6785a2e... is provided.

The tweet includes a screenshot of the Mempool.space transaction page for **Transaction c007dc731e27d53521249bdbd4f0b2f472bcb911a...**. The transaction details are as follows:

First seen	9 hours ago	Fee	12,636,000 sat \$7,842
Features	SegWit Taproot RBF	Fee rate	25,867 sat/vb

The **RBF History** section shows a timeline of replacements:

Time	Fee (sat/vb)
0:10	569
	653
	961
	975
	1,476

The tweet was posted at 11:47 PM · Mar 19, 2024 and has 20.4K Views. It has 13 replies, 33 retweets, 137 likes, and 34 bookmarks.



“mempool sniping”

Lesson: if you sign a transaction with
SIGHASH_ANYONECANPAY...

... anyone can change its inputs



👉 don't design stuff like this

👉 we can't fix this in transaction relay

mononaut (tx/acc) @mononautical

Mempool sniping is going wild.

This ordinal trade was broadcast at the listed price of 0.069 BTC, and within a few minutes it had been replaced at least 14 times by competing buyers, bidding the fees up to 0.126 BTC (25,867 s/vb)


[mempool.space/tx/c04a6785a2e...](#)

Transaction

c007dc731e27d53521249bdbd4f0b2f472bcb911a... f7c630e3c019 Removed

First seen	9 hours ago	Fee	12,636,000 sat \$7,842
Features	SegWit Taproot RBF	Fee rate	25,867 sat/vb

RBF History








The RBF History timeline shows a sequence of transactions being replaced. It starts with a transaction at 0:10 with a fee of 569 sat/vb. This is replaced by a transaction with a fee of 653 sat/vb, which is then replaced by a transaction with a fee of 961 sat/vb, and so on, up to a final transaction with a fee of 1,476 sat/vb. The timeline is marked with time intervals: 20 seconds, 5 seconds, 14 seconds, and 8 seconds.

569	653	961	975	1,476
0:10				

11:47 PM · Mar 19, 2024 · 20.4K Views


13 33 137 34

We can test censorship-resistance

- Add 1 honest peer + n “evil” peers programmed with possible actions:
 - forgetting 
 - stalling 
 - flooding 
 - malleating txs 
 - disconnecting 
- Simulate random combinations of messages and assert that we always receive the tx



so if the network looks like this, the transaction can still go through



transaction relay vs
computation



computation has to be limited

our operating environment = anonymous entities on the internet
connect to you and send you data to process

- 👉 blocks are naturally DoS-resistant (must have PoW!)
- 👉 transactions can be computationally expensive to handle:
 - signature verification
 - mempool package updates
 - etc

signature verification

quadratic sighashing and the 1MB “Megatransaction”

Rusty Russell's Quiet Corner of The Internet

About

The Megatransaction: Why Does It Take 25 Seconds?

Jul 8, 2015

Last night f2pool mined a 1MB block containing a single 1MB transaction. This scooped up some of the spam which has been going to various weakly-passworded “brainwallets”, gaining them 0.5569 bitcoins (on top of the normal 25 BTC subsidy). You can see the megatransaction on [blockchain.info](#).


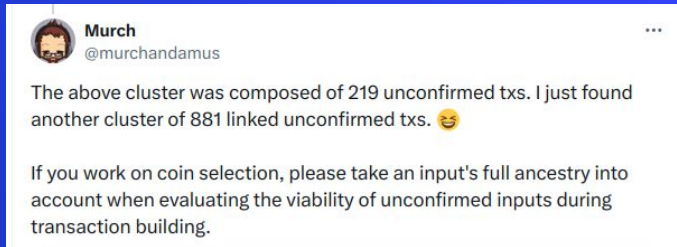
It was widely reported to take about 25 seconds for bitcoin core to process this block: this is far worse than my “2 seconds per MB” result in my last post, which was considered a pretty bad case. Let's look at why.

👉 only validate transactions of limited size and sigops,
only allow up to 15-of-15 legacy multisig

mempool package updates

insertion and removal require updating ancestors and descendants

👉 limit the number and size of ancestor and descendant packages



👉 ... limit the cluster size too


Transaction Relay / Mempool Policy

Bitcoin Core rejects some consensus-valid transactions

- 📌 **decentralization**: minimize memory, bandwidth, computation
- 📌 **security**: prevent DoS, specific bugs not fixed in consensus
- 📌 **upgrade hooks**: don't use future things before they are defined (avoid collisions and confiscation)
- 📌 **best practices aka paternalism**: discouraging network resource-intensive practices



default policy is *like* a transaction relay interface, but it's optional and varies



Wait, but isn't this CENSORSHIP?



They're all tradeoffs

- accuracy of assessing incentive compatibility / computational complexity
- allowing complex topologies / fixing pinning (censorship) problems
- relaying all consensus-valid transactions / safely updating consensus later

👉 to design for censorship-resistance (and other goals), we have to reject some

👉 the point is to agree on an interface that is both useful for users and acceptable for node operators, not stop users from broadcasting transactions

mempool policies can differ, but it's best for caches to be hot




Practice question from twitter

- In the past, some rules were added to discourage practices proactively
 - uneconomical (“dust”) outputs
 - large (nonstandard) scriptPubKeys

Q: Why can't Bitcoin Core just filter {inscriptions, ordinals, brc-20s} too?
#FixTheFilters

“Just make Bitcoin Core filter these transactions”



miners are not obligated to
run software that throws
away fee revenue



this doesn't work on
transactions that already
bypass policy rules



if you discard transactions
that are likely to be mined,
your cache is useless

Summary

- ✂️ Design goals like “censorship resistance” and “decentralization” aren’t easy to achieve
- ✂️ They do create interesting technical puzzles
- ✂️ All the code is open source - that’s important
- ✂️ YOU can see for yourself how it works



Go deeper!



Subscribe to the Optech Newsletter

bitcoinops.org

Attend a Chaincode Learning Seminar

chaincode.gitbook.io/seminars

Join Bitcoin Core PR Review Club

bitcoincore.reviews

Contribute to the code

github.com/bitcoin/bitcoin

Thanks!

Any questions?

@glozow

